Strategic Roadmap for High-Resolution Electrical Load Forecasting and Grid Stability: The GRIDSHIELD Case Study

The execution of high-frequency electrical load forecasting within the pressurized environment of a business hackathon requires a rigorous synthesis of technical data science, behavioral economics, and deep regulatory knowledge. The GRIDSHIELD Case Study, presented as Case 2 in the Decode X-2026 Finale Round, places participants in the role of strategic advisors for a Mumbai-based power distribution company (DISCOM). The objective is to develop a two-day-ahead load forecasting model with 15-minute resolution that is capable of maintaining grid stability and economic efficiency amidst structural shocks and behavioral regime shifts.[1] This roadmap provides an exhaustive guide for teams to navigate the 22-hour competition window, moving from baseline data analysis to the final executive board presentation, while addressing the unique complexities of Mumbai's energy landscape and the aftermath of the COVID-19 pandemic.[1]

The Competitive Landscape of Decode X-2026

Success in the Decode X-2026 hackathon is predicated on more than technical accuracy; it requires a holistic understanding of how data translates into strategic decisions. The competition is structured to simulate real-world volatility, where baseline models are frequently upended by sudden policy changes or environmental events.[1] The 22-hour timeline is a critical constraint that necessitates a clear division of labor among team members, ensuring that while the technical pipeline is being constructed, the strategic narrative and regulatory context are concurrently developed.

Competition Milestone	Schedule Detail	Critical Deliverables and Objectives
Commencement	28 February, 11:00 AM	Baseline data ingestion and exploratory analysis.[1]
Clarification Period	11:00 AM – 1:00 PM	Addressing business interpretations of Case 2 variables.[1]
Strategic Development	1:00 PM – 7:00 PM	Building the "Stage 1" baseline predictive model.[1]

Stage 2 Release	28 February, 07:00 PM	Introduction of the "Structural Shock" or regime shift.[1]
Interim Submission	28 February, 11:00 PM	5-slide shock diagnostic and recalibration summary.[1]
Stage 3 Release	01 March, 01:00 AM	Confidential board directive and binding constraints.[1]
Final Submission	01 March, 09:00 AM	15-slide presentation, technical appendix, and memo.[1]
Jury Evaluation	Post-Submission	20-minute presentation and 10-minute Q&A for finalists.[1]

The evaluation criteria reflect the multidisciplinary nature of the task, with Analytical Depth and Decision Logic each accounting for 25% of the total score. The ability to re-optimize under the Stage 2 shock carries a 20% weight, emphasizing the need for a model architecture that is not just precise but also adaptable.[1] Business Communication and Systems Thinking round out the evaluation, requiring teams to articulate how their forecasting model interacts with the broader Maharashtra grid ecosystem.[1]

Structural Analysis of Mumbai's Energy Ecosystem

Mumbai presents a uniquely challenging energy profile due to its status as India's financial capital, its high population density, and its humid subtropical climate. The load data provided in Case 2, covering 15-minute intervals, captures the rhythmic pulse of the city’s economic and social life.[1] To build an effective forecast, one must recognize that the load curve is not merely a sequence of numbers but a reflection of the synchronous behaviors of millions of residents and thousands of commercial entities.[1]

High-Resolution Load Profiling

The 15-minute resolution of the Electric_Load_Data_Train.csv file allows for the identification of high-frequency patterns that are essential for short-term load forecasting (STLF).[1] These intervals align with the time-block concept used in energy accounting and deviation settlement.[2] Analysis of the training data reveals typical daily cycles: a late-night baseload, a sharp morning ramp-up as commercial activities commence, sustained high demand during business hours, and a secondary peak in the evening as residential consumption surges.[1]

Load Variable	Granularity	Historical Context
Target (LOAD)	15-minute intervals	kW consumption captured from 2013 to 2021.[1]
Temporal Range	8 years	Covers pre-COVID, COVID shock, and recovery.[1]
Forecasting Horizon	48 hours	2-day-ahead scheduling for SLDC compliance.[1]

The case speaker identifies "layers of patterns" that must be unpeeled. This implies that the load is a composite of different consumer segments: residential, commercial, and industrial. Each segment reacts differently to temporal and environmental variables. For example, commercial load is highly predictable on weekdays but drops significantly on Sundays, while residential load may peak during late evening hours and remains relatively higher during weekends.[1]

Environmental Determinants and the COOL_FACTOR

The External_Factor_Data_Train.csv provides environmental variables that are the primary drivers of seasonal load variation in Mumbai.[1] In a subtropical climate, the relationship between temperature and load is nonlinear, largely governed by the usage of air conditioning (AC) systems.

The dataset introduces a unique engineered feature: COOL_FACTOR. Analysis of the data indicates that this factor operates on a threshold-based logic, serving as a proxy for the cooling demand generated by AC units.[1] The COOL_FACTOR remains at zero when ambient temperatures are below a comfortable threshold (approximately $27^{\circ}C$).[1] Once the temperature exceeds this point, the COOL_FACTOR begins to track the thermal intensity, indicating an immediate and widespread activation of cooling infrastructure across the city.

DateTime Sample	Actual Temperature (◦C)	COOL_FACTOR Value	Demand State
01APR2013:10:30:00	26.68	0.00	Threshold not reached.[1]
01APR2013:10:45:0	27.15	0.15	Initial AC

0			activation.[1]
01APR2013:11:00:00	27.32	0.32	Rising cooling demand.[1]
01APR2013:15:00:00	32.89	5.89	Peak thermal load.[1]

The interaction between temperature and humidity is equally critical. Mumbai's high humidity prevents natural cooling, leading to a higher "heat index" and sustained AC usage even after sunset.[1] Forecasting models must utilize these variables not just as standalone features but as interaction terms to capture the perceived heat that drives consumer behavior. Furthermore, the data indicates that every 15-minute interval is recorded with zero ACT_RAIN for long stretches, suggesting that while rain can provide sudden cooling, it may also lead to "weather holidays" where business activity slows down significantly.[1]

The Taxonomy of Mumbai Festivals and Event-Driven Load

Mumbai's electricity demand is fundamentally tethered to its cultural calendar. The Events_Data.csv provides a detailed list of festivals and holidays that create significant departures from the baseline load profile.[1] These events must be treated as critical categorical features in any forecasting model.

Single-Day vs. Multi-Day Event Impact

Mumbai's calendar is unique because of the prevalence of multi-day religious observations. Unlike a standard national holiday like Republic Day, festivals such as Ganesh Chaturthi or Navratri span 7 to 10 days.[1] During these periods, the city's residential load increases due to evening celebrations, decorative lighting, and social gatherings, while industrial and commercial sectors may operate at reduced capacity.

Event Type	**Event Examples**	**Load Profile Characteristics**
National Holidays	Independence Day, Gandhi Jayanti	Drastic reduction in commercial/industrial load.[1]
Religious Single-Day	Eid, Christmas, Makar	Shifted morning peak;

	Sankranti	residential evening peak.[1]
Multi-Day Festivals	Ganesh Chaturthi, Navratri	Sustained baseline shift over several days.[1]
Weather-Induced	Very Heavy Rain	Treated as a holiday in binary indicators (1).[1]
Administrative	Election Day, Bandhs	Sunday-like load profiles on weekdays.[1]

A critical insight from the events data is the classification of "Very Heavy Rain" as a holiday (Holiday_Ind = 1).[1] This reflects the operational reality of Mumbai, where extreme monsoonal events often lead to a total cessation of public transport and business activity. For a forecaster, this means that the model must not only predict the cooling effect of rain but also the drastic "holiday effect" it induces in commercial sectors. Failing to account for this dual impact would lead to significant over-forecasting during the monsoon season.

The Behavioral Dynamics of Ganesh Chaturthi

Ganesh Chaturthi is the defining cultural event of Mumbai, and its impact on the grid is complex. The data distinguishes between "Day 1," "Ganesh Festival Day 2-10," and "Ganesh Visarjan".[1] During the final day (Visarjan), mass processions occur across the city, leading to widespread closures of shops and offices, resulting in a load curve that may even be lower than a typical Sunday. Models must incorporate "event-lag" features that track the progression of these multi-day festivals to capture the cumulative behavioral shifts of the population.

The COVID-19 Structural Shock and Behavioral Regime Shifts

The inclusion of 2020 data within the training set introduces a "structural shock" that is central to the GRIDSHIELD case.[1] The COVID-19 pandemic did not merely cause a temporary drop in demand; it reconfigured the social and economic structures of the city, leading to permanent shifts in the energy profile.

The "Weekdays Resemble Sundays" Phenomenon

During the stringent lockdown phases in 2020, the Indian power system experienced unprecedented changes in the shape of its load curve. Regression analyses conducted by energy researchers found that nearly 73% of the total drop in demand was attributed to the industrial sector, while 27% was due to the commercial sector.[3] Residential demand, conversely, remained stable or increased due to "stay-at-home" mandates.[4]

This resulted in a "flattening" of the daily load curve. The sharp morning peak, typically driven by office and factory operations, largely disappeared. Stay-at-home populations reached 40% in Mumbai, causing a shift in consumption toward the residential sphere.[5] The resultant curve shape prompted industry analysts to claim that "weekdays resemble prolonged Sundays," characterized by a lack of morning ramp-up and a more distributed afternoon load.[3]

Segment	Lockdown Impact	Post-Shock Baseline
Industrial	73% of total demand drop.[3]	Slower recovery; linked to supply chains.[4]
Commercial	27% of total demand drop.[3]	Permanent shift toward remote/digital work.[6]
Residential	Increased share of total load.[4]	New baseline due to hybrid work models.[1]

Permanent Post-Pandemic Regime Shifts

The speaker in the provided transcript emphasizes that the shifts occurring after COVID-19 are permanent.[1] This "new normal" is driven by a digitized service economy and the adoption of hybrid work structures.[6] For the modeler, this means that data from 2013-2019 may no longer accurately represent the baseline behavior of the 2021 grid.

Teams must implement a regime-switching framework to handle this shock. In Stage 2 of the hackathon, it is highly likely that an "out of time sample" from the post-lockdown era will be introduced to test how models adapt to this new behavioral baseline.[1] Effective strategies include using the "Stringency Index" as a feature to capture the intensity of behavioral restrictions or applying a higher weight to more recent post-2020 data points.[5]

Technical Modeling Roadmap for High-Granularity Forecasting

Developing a model for a 15-minute resolution, 48-hour forecasting horizon is an exercise in managing both long-term seasonality and short-term volatility. The provided training set contains approximately 280,000 observations, making it ideal for deep learning and advanced ensemble techniques.[1]

Feature Engineering and Data Preprocessing

The foundation of the forecast is a robust feature set that captures the "layers of patterns."

Beyond basic weather variables and event dummies, several engineered features are mandatory for a winning solution.

- **Temporal Decomposition:** Extracting the hour of day (0-23), day of the week (0-6), and month of the year (1-12). A binary "Weekend" flag and a "Day of Week" sine/cosine transformation should be used to capture the cyclic nature of time.
- **Auto-Regressive Lags:** Load values from exactly 24 hours ago, 48 hours ago, and 7 days ago. For a 48-hour ahead forecast, the model cannot use the "last 15 minutes" of load as a predictor; it must rely on lags that are at least 48 hours old.[1]
- **Rolling Weather Statistics:** Moving averages of temperature and humidity over the preceding 6 to 12 hours. This captures the "thermal inertia" of Mumbai's buildings, where sustained heat leads to higher AC demand regardless of a slight dip in instantaneous temperature.
- **Event Proximity Indicators:** Creating "Lead" and "Lag" dummies for major festivals. For example, if today is the day before Diwali, the commercial load might increase as people shop, whereas the day after Diwali may see a sharp drop.[9]

Selection of Predictive Algorithms

Given the complex nonlinear interactions between COOL_FACTOR, religious events, and structural shocks, simple linear models will underperform. The research identifies three primary architectures that dominate electricity load forecasting.

1. **Ensemble Tree Models (XGBoost / LightGBM):** These are highly effective at handling tabular data with categorical variables (festivals) and threshold features (COOL_FACTOR).[10] They are robust to outliers and can model the sharp discontinuities introduced by monsoon-related holidays.
2. **Recurrent Neural Networks (LSTM / GRU):** Long Short-Term Memory networks are specifically designed for sequential data.[12] They are exceptionally good at learning long-term seasonal dependencies and the evolving behavioral patterns of a post-COVID population. A GRU model often provides comparable accuracy to LSTM with lower computational overhead in 15-minute resolutions.[12]
3. **Quantile Regression (QR):** Traditional models minimize Mean Squared Error (MSE), which assumes a symmetric penalty for over- and under-forecasting. In a business context, under-forecasting is far more expensive.[13] QR allows the DISCOM to forecast a specific percentile (e.g., the 80th percentile) to ensure a safety margin against blackouts.[15]

Model Class	Strength	Best Application in GRIDSHIELD
Gradient Boosting	Robustness to non-linear	Modeling the COOL_FACTOR activation

	thresholds	point.[1]
LSTM / GRU	Captures long-term cyclic dependencies	Mapping the 8-year evolution of Mumbai's load.[12]
Quantile Regression	Asymmetric risk management	Minimizing high-cost frequency-linked DSM penalties.[16]

The Economics of Grid Stability: ABT and DSM Regulations

A fundamental requirement of the GRIDSHIELD case is understanding the regulatory framework that governs the Mumbai power market. The Maharashtra Availability Based Tariff (ABT) and the Deviation Settlement Mechanism (DSM) provide the commercial incentives that drive forecasting precision.[1]

The Mechanism of Deviation Penalties

The ABT mechanism requires DISCOMs to submit their power drawal schedules to the State Load Dispatch Centre (SLDC) for every 15-minute block of the coming day.[1] A "Deviation" is defined as the difference between the actual drawal and the scheduled drawal.[17]

The "Normal Rate" of deviation charges is typically linked to the Market Clearing Price on power exchanges like I-DAM (Integrated Day-Ahead Market) or RTM (Real-Time Market).[17] However, the actual penalty is frequency-linked. When the grid frequency drops below 49.85 Hz—indicating that demand is exceeding supply—the cost of over-drawal (under-forecasting) increases to 150% or even 200% of the normal rate to discourage actions that could lead to a grid collapse.[17]

Grid Condition	**Forecasting State**	**Commercial Consequence**
Frequency < 49.85 Hz	Over-drawal (Under-forecast)	Penalties of 150-200% of Normal Rate.[17]
Frequency > 50.05 Hz	Under-drawal (Over-forecast)	Zero or negative receivable; revenue loss.[17]

Volume > 10% or 30 MW	High Deviation	Additional surcharges and regulatory scrutiny.[18]

For a DISCOM, under-forecasting during peak hours is a catastrophic economic risk. It results in both high DSM penalties and potential blackouts that damage infrastructure and consumer trust.[13] Over-forecasting, while less risky for grid stability, leads to the purchase of excess power that may be sold back to the pool at a significant loss, especially if the grid is already over-supplied (high frequency).[17]

Asymmetric Optimization Strategies

To win a business hackathon, the model must be optimized for "Forecast Value" rather than just "Forecast Accuracy".[23] This means implementing a piecewise loss function that penalizes under-predictions more heavily than over-predictions.

Mathematically, the objective shifts from minimizing $\sum(y - \hat{y})^2$ to a cost-oriented function:

$$L = \begin{cases} k_1 \cdot (y - \hat{y}) & \text{if } \hat{y} > y \text{ (Over-prediction)} \\ k_2 \cdot (y - \hat{y}) & \text{if } \hat{y} \leq y \text{ (Under-prediction)} \end{cases}$$

where $k_2 >> k_1$, reflecting the higher cost of over-drawal penalties compared to surplus procurement costs.[22] This strategy, implemented through Quantile Regression, can reduce daily economic costs by 42% to 57% depending on the actual cost ratio of the two types of errors.[13]

Strategic Roadmap for Stages 2 and 3

The evolution of Case 2 through Stages 2 and 3 is designed to test the robustness of the team's thinking process. Participants must be prepared to pivot their technical approach in response to sudden case directives.[1]

Navigating Stage 2: The Structural Shock

Stage 2 introduces a "regime shift" at 7:00 PM on the first day.[1] Based on the speaker's hints, this could involve the release of an "out of time sample" where the load pattern has radically diverged from the 2013-2019 baseline.[1]

Potential Stage 2 Scenarios:

- **Post-COVID Permanence:** The introduction of 2021-2022 data that shows a permanent flattening of the load curve due to hybrid work policies.[1]
- **Regulatory Tightening:** A narrowing of the DSM tolerance band (e.g., from $\pm 10\%$ down

to $\pm 5\%$), making deviations significantly more expensive.[20]
- **EV Mass Adoption:** A shock scenario where residential EV charging peaks are superimposed on existing evening peaks, creating "local overloading" risks for transformers.[24]

The interim submission must quantify the impact of this shock. Teams should compare their baseline model's performance on the new sample versus its performance on the historical training set. The "Shock Diagnosis" should clearly articulate whether the change is one of magnitude (demand level) or shape (load profile) and what trade-offs are required to maintain grid stability under these new conditions.[1]

Mastering Stage 3: The Board Directive

At 1:00 AM, the "Final Board Directive" introduces a confidential binding constraint.[1] This represents a revised business objective that must be incorporated into the final model and strategy.

Strategic adaptation for Stage 3:

- **Constraint Integration:** If the board mandates a "Zero Under-Forecast" policy during evening peaks, the team must switch from a mean forecast to a high-quantile forecast (e.g., 95th percentile) for those specific blocks, even if it increases the average MAPE.
- **Cost vs. Risk Optimization:** Using Mixed Integer Linear Programming (MILP) to find the optimal scheduling mix given the forecast and the new constraint.[26] For instance, the directive might require minimizing total procurement costs while keeping the probability of a 30 MW deviation below 1%.
- **Systems Thinking:** The final recommendation must explain how the model serves as a "GRIDSHIELD," protecting the transformer and cable infrastructure from unpredicted real-time congestion.[25]

The Final Deliverable Suite: Professional Communication

The conclusion of the 22-hour cycle is the submission of a tripartite deliverable package: the Executive Board Presentation, the Technical Appendix, and the One-Page Decision Memo.[1]

Executive Board Presentation (Max 15 Slides)

The presentation must be framed for a non-technical audience of utility executives. It should tell a story that begins with the baseline diagnostic and concludes with a resilient strategy for the future.[1]

1. **Problem Framing:** Identifying the 15-minute load forecast as the critical link between

DISCOM profitability and Mumbai's energy security.[1]

2. **Baseline Diagnosis:** Visualizing the pre-shock load curve and its correlation with COOL_FACTOR and religious events.[1]
3. **Shock Impact Analysis:** Showing the "flattening" effect of the 2020 lockdowns and quantifying the error increase if models are not recalibrated.[3]
4. **Optimization Logic:** Explaining the switch to asymmetric loss functions and quantile regression to mitigate frequency-linked DSM penalties.[13]
5. **Final Recommendation:** A clear directive for scheduling that balances cost and grid risk, accounting for the binding Stage 3 constraints.[1]

Technical Appendix (Max 10 Pages)

The technical appendix provides the evidentiary foundation for the presentation. It should be a professional document indistinguishable from a consultant's technical report.[1]

- **Model Specifications:** Detailed architecture of the XGBoost or LSTM models, including hyperparameter tuning results (e.g., learning rates, dropout, layer depth).[12]
- **Assumption Validation:** A critical analysis of why certain variables (like ACT_RAIN) were included or excluded and how the threshold logic for COOL_FACTOR was determined.[1]
- **Sensitivity Analysis:** Demonstrating how the model responds to a $2^\circ C$ temperature rise or a one-day shift in a festival date. This proves that the model has captured the "layers of patterns" rather than just memorized the training data.[1]

One-Page Decision Memo

The memo is the most critical document for a high-stakes business decision. It must summarize the "fundamental change" that occurred during the simulation (the shock), the specific trade-offs the team accepted (e.g., higher procurement cost for lower deviation risk), and the final recommended action.[1]

Conclusion: Strategic Synthesis for GRIDSHIELD Victory

The GRIDSHIELD case study is an rigorous test of an analytical team's ability to synthesize disparate data streams into a cohesive grid management strategy. In the volatile 22-hour landscape of Decode X-2026, the winners will be those who recognize that the 15-minute load profile of Mumbai is a social document, reflecting the interplay of climate, culture, and structural economic shocks.[1]

By leveraging threshold-based environmental features like the COOL_FACTOR and recognizing the permanent behavioral regime shifts induced by the COVID-19 pandemic, teams can build a technical pipeline that is both accurate and adaptable.[3] Integrating this technical prowess with

a sophisticated understanding of the Maharashtra DSM regulatory framework—specifically the asymmetric costs of under-forecasting—allows for the creation of a "GRIDSHIELD" that protects both the financial viability of the DISCOM and the stability of the urban grid.[17]

The Roadmap to Success:

- **Exploiting High Resolution:** Use the 15-minute granularity to capture the synchronous behavioral spikes of a megacity.[1]
- **Threshold Awareness:** Engineers features that trigger at specific climate points, mimicking real-world AC activation logic.[1]
- **Behavioral Regime Mapping:** Model the 2020 shock as a permanent shift in baseline demand rather than a transient anomaly.[1]
- **Economic Risk Optimization:** Optimize for the asymmetric penalty landscape of frequency-linked deviation charges.[13]
- **Evolutionary Agility:** Maintain a modular model architecture that can be quickly re-optimized for the Stage 3 binding board directives.[1]

In the final analysis, the GRIDSHIELD challenge is won by the team that provides the most robust "thinking process," demonstrating that they can navigate the complex trade-offs between accuracy, cost, and reliability in one of the world's most dynamic energy markets.[1]

Works cited

1. Case 02 - Stage Guidelines.pdf
2. Summary of HT Tariff for FY 2024-25, effective from 1 April, 2024 8.13 Recording of Maximum Demand MSEDCL's Submis - The Leading Solar Magazine In India, accessed on February 28, 2026, https://www.eqmagpro.com/wp-content/uploads/2020/03/Order-322-of-2019_compressed-501-625.pdf
3. Analysing the impact of lockdown due to the COVID-19 pandemic on the Indian electricity sector - PMC, accessed on February 28, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC9550349/
4. Change in the shape of the normalised average load curve of Mumbai in successive lockdown phases. - ResearchGate, accessed on February 28, 2026, https://www.researchgate.net/figure/Change-in-the-shape-of-the-normalised-average-load-curve-of-Mumbai-in-successive-lockdown_fig14_359340233
5. Sensing global changes in local patterns of energy consumption in cities during the early stages of the COVID-19 pandemic - PMC, accessed on February 28, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC9212780/
6. GDP Rebase 2022–23: What Changes in India's Growth Calculations - Outlook Business, accessed on February 28, 2026, https://www.outlookbusiness.com/economy-and-policy/gdp-rebase-202223-what-changes-in-indias-growth-calculations
7. The impact of COVID-19 measures on intraday electricity load curves in the

European Union: A panel approach - PMC, accessed on February 28, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC9484106/
8. Empirical Analysis of the Impact of COVID-19 Social Distancing on Residential Electricity Consumption Based on Demographic Characteristics and Load Shape - MDPI, accessed on February 28, 2026, https://www.mdpi.com/1996-1073/14/22/7523
9. Household Electricity Load Forecasting for Demand Side Management Applications - UT Student Theses, accessed on February 28, 2026, https://essay.utwente.nl/fileshare/file/94510/Kok_MA_EEMCS.pdf
10. Electric Load Forecasting Methods: A Comparison - Yes Energy, accessed on February 28, 2026, https://www.yesenergy.com/blog/electric-load-forecasting-methods
11. GridShield: A Cyber-Secure AI Framework for Real-Time Detection of Electricity Theft Based on Daily Consumption Behavior - ResearchGate, accessed on February 28, 2026, https://www.researchgate.net/publication/397357070_GridShield_A_Cyber-Secure_AI_Framework_for_Real-Time_Detection_of_Electricity_Theft_Based_on_Daily_Consumption_Behavior
12. Electrical Load Forecasting Using LSTM, GRU, and RNN Algorithms - MDPI, accessed on February 28, 2026, https://www.mdpi.com/1996-1073/16/5/2283
13. An Asymmetric Loss with Anomaly Detection LSTM Framework for Power Consumption Prediction - arXiv.org, accessed on February 28, 2026, https://arxiv.org/pdf/2302.10889
14. Support vector regression with asymmetric loss for optimal electric load forecasting, accessed on February 28, 2026, https://ideas.repec.org/a/eee/energy/v223y2021ics0360544221002188.html
15. Quantile Regression for Probabilistic Electricity Price Forecasting in the U.K. Electricity Market - IEEE Xplore, accessed on February 28, 2026, https://ieeexplore.ieee.org/iel8/6287639/10820123/10838567.pdf
16. Quantile Regression Model for Peak Load Demand Forecasting with Approximation by Triangular Distribution to Avoid Blackouts# - EconJournals.com, accessed on February 28, 2026, https://www.econjournals.com/index.php/ijeep/article/download/6742/3933/17209
17. Maharashtra Issues Draft Deviation Settlement Mechanism Regulations - Mercom India, accessed on February 28, 2026, https://www.mercomindia.com/maharashtra-issues-draft-deviation-settlement-mechanism-regulations
18. draft maharashtra electricity regulatory commission (deviation, settlement mechanism and related matters) regulations, 2025, accessed on February 28, 2026, https://avantiscdnprodstorage.blob.core.windows.net/legalupdatedocs/45785/MERC-issued-a-public-notice-inviting-comments-on-the-Draft-MERC-Deviation-Settlement-Mechanism-and-Related-Matters-Regulations-2025-AUG202025.pdf
19. Bid Areas and Deviation Charges Overview | PDF - Scribd, accessed on February 28, 2026,

https://www.scribd.com/document/905268486/bidArea-calculation-matrix
20. Rigid Deviation Settlement Limits Could Result in Higher Renewable Tariffs - Mercom India, accessed on February 28, 2026, https://www.mercomindia.com/deviation-settlement-higher-tariffs
21. The Billion-Dollar Costs of Forecasting Electricity Demand - RMI, accessed on February 28, 2026, https://rmi.org/billion-dollar-costs-forecasting-electricity-demand/
22. Personalized Federated Learning with Cost-Oriented Load Forecasting for Home Energy Management Systems - UPCommons, accessed on February 28, 2026, https://upcommons.upc.edu/bitstreams/ddbf9cf4-3ce2-478d-8a52-71428d1a77c0/download
23. Assessing and Comparing Short Term Load Forecasting Performance - MDPI, accessed on February 28, 2026, https://www.mdpi.com/1996-1073/13/8/2054
24. GridShield—Optimizing the Use of Grid Capacity during Increased EV Adoption, accessed on February 28, 2026, https://www.researchgate.net/publication/369111824_GridShield-Optimizing_the_Use_of_Grid_Capacity_during_Increased_EV_Adoption
25. GridShield—Optimizing the Use of Grid Capacity during Increased EV Adoption - MDPI, accessed on February 28, 2026, https://www.mdpi.com/2032-6653/14/3/68
26. SHORT TERM LOAD FORECASTING USING QUANTILE REGRESSION WITH AN APPLICATION TO THE UNIT COMMITMENT PROBLEM Moshoko Emily Lebotsa 1 - Univen, accessed on February 28, 2026, https://univendspace.univen.ac.za/bitstreams/9b34d7ad-7f11-497e-bf4d-83b6afba26e2/download
27. Charging electric vehicles, eating pizzas, and safeguarding the grid in 2025 - IET Digital Library, accessed on February 28, 2026, https://digital-library.theiet.org/doi/pdf/10.1049/icp.2025.1605?download=true